

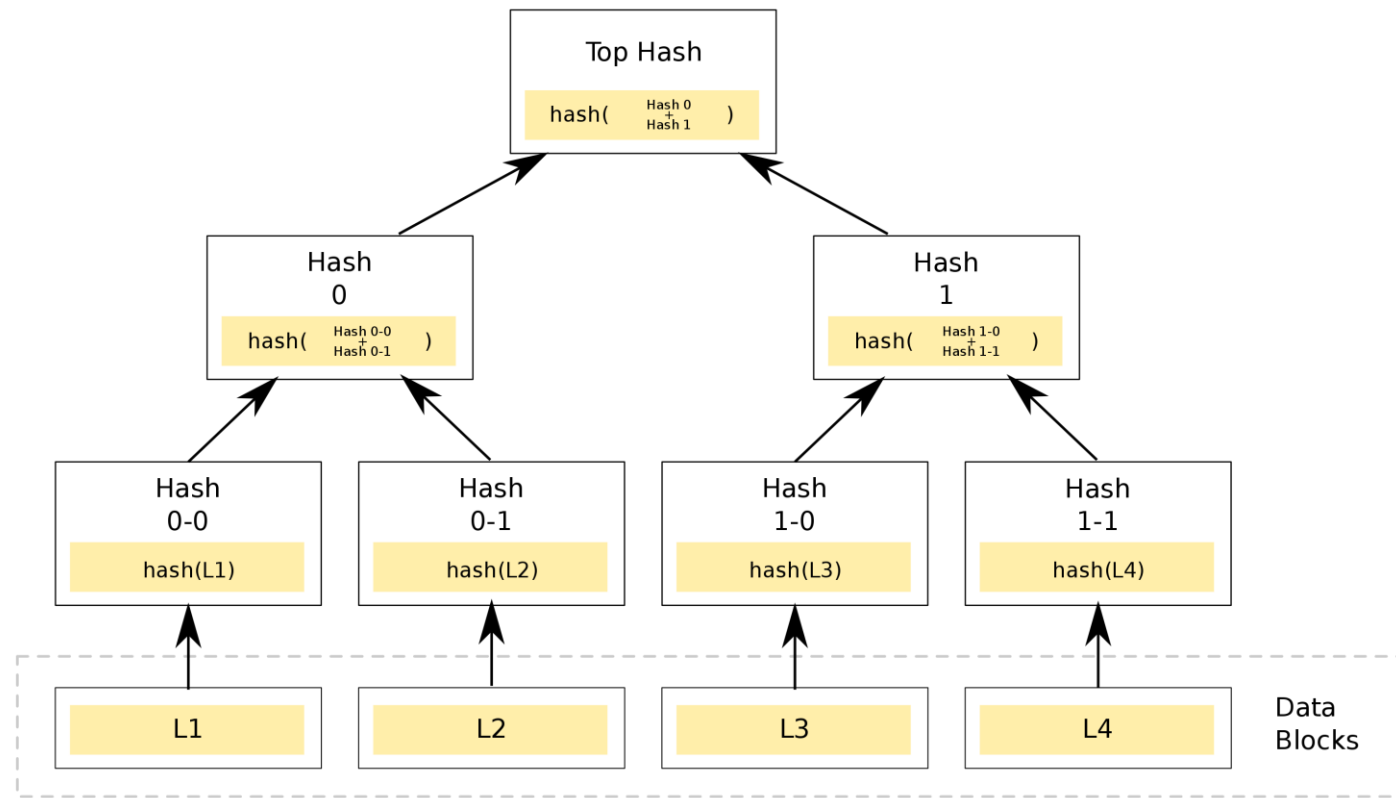
MerkleTrees



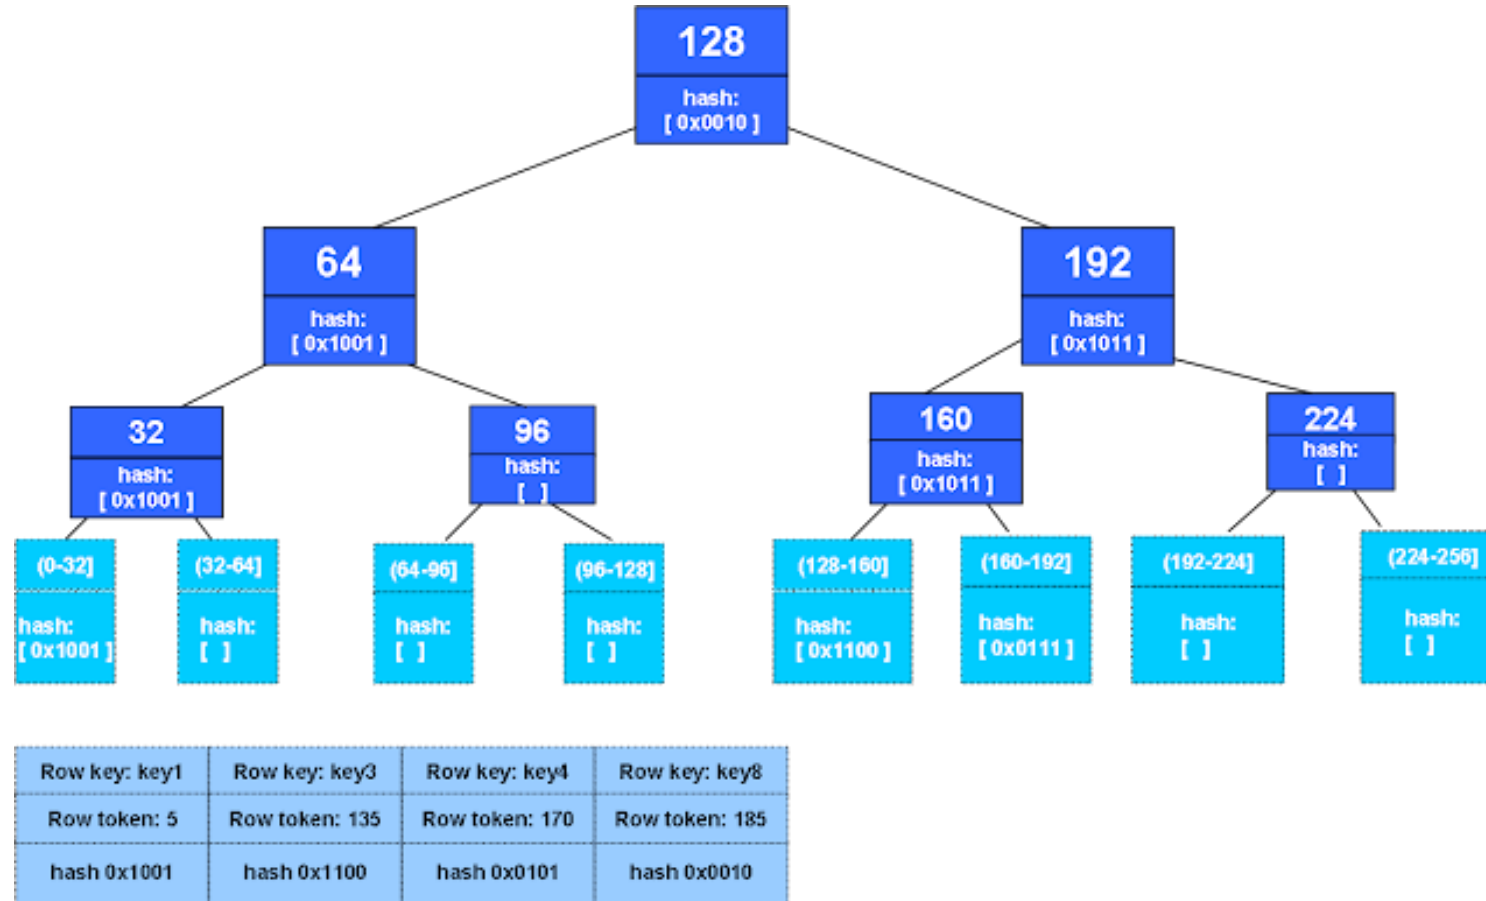
# How would you synchronize 2 databases?

- Both have millions of rows
- Crappy ingestion resulted in many rows missing
- Crappy internet makes batch dump too costly

# MerkleTree example



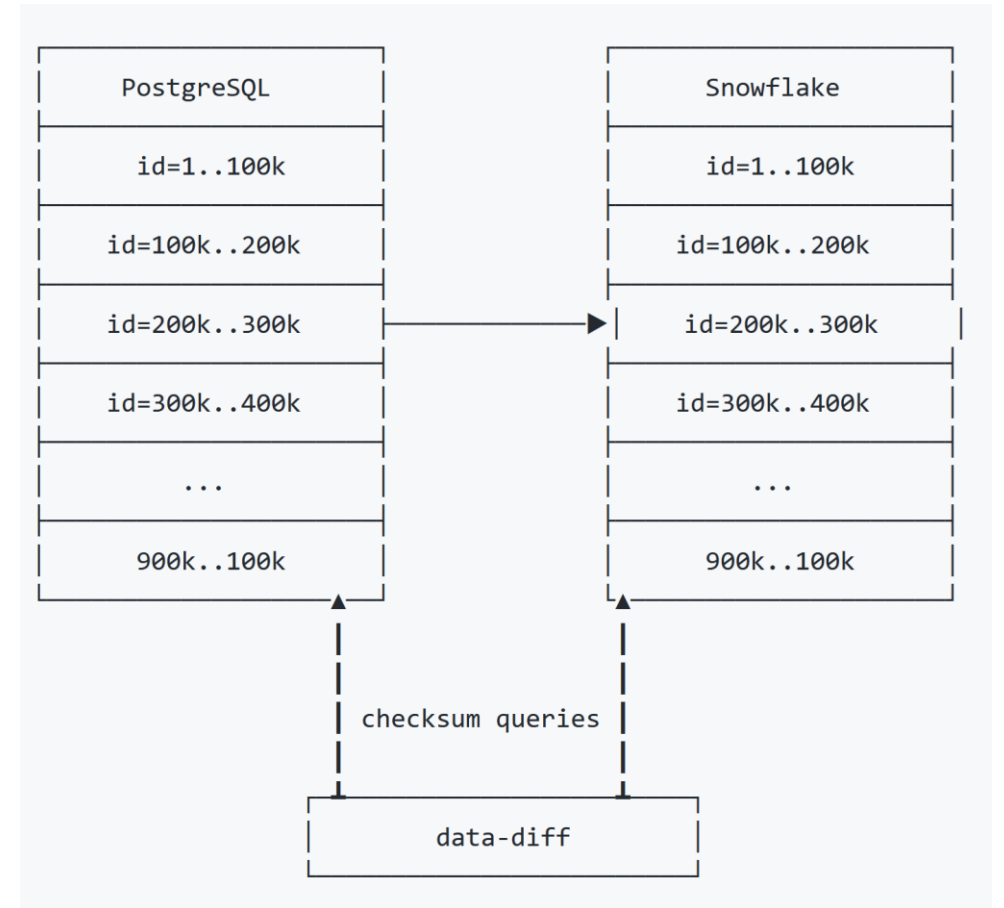
# MerkleTree Cassandra



# Data-Diff

## Poor-mans merkletree

1. min/max on id column
2. divide into X buckets
3. hash each bucket
4. compare hashes
5. repeat



# MerkleTree, hackaton

- [https://github.com/NielsZeilemaker/merkle\\_tree\\_tutorial](https://github.com/NielsZeilemaker/merkle_tree_tutorial)